LEXIGRID

Document Extraction and ISO Standard Review

Version 5.54.15 | Reviewed 2026-11-11

Shahryar Jahangir, Aniza Inc.

Procedures, examples, verification and operational limits. Complete commissioning on your deployment before serving customer traffic. Sample values are fictional; original keys and data must be preserved.

Contents

Reviewed document import **3**

What to review . . . 3

Select all and complete the import . . . 3

Import behavior . . . 4

Scope and limitations . . . 4

ISO/IEC 27001:2022 PDFs . . . 4

Import requirements when no framework exists . . . 5

Large PDF documents and page ranges . . . 5

NIST SP 800-53 Revision 5 catalogue extraction . . . 5

PDF layout padding and extraction diagnostics . . . 6

Choose the record type and verify the destination . . . 6

Default accountable owner and duplicate control imports . . . 6

CJIS Security Policy version 6.1 . . . 7

Requirements navigation . . . 7

Reviewed document import

Use **Document Extraction** to import controls, frameworks, policies, or requirements from text-based PDF, DOCX and UTF-8 TXT files. Select the record type before extraction. Changing the file or type clears the old preview. Extraction does not save records and all approval checkboxes start unchecked.

What to review

- **Controls:** review code, title, description and owner. Numbered/control-code headings and recognized table columns are supported, including common references such as AC-2, AC-2(1), A.5.15, CC6.1 and GV.OC-01. Descriptions need not contain "must" or "shall." Numbered actions under a control remain with that control. If a single statement lacks a separate title, the title stays blank for review. The optional default owner is entered by the user and fills only missing source owners.

- **Frameworks:** review name and revision. Recognized table headers and explicit document labels such as `Framework:` and `Version:` are supported. Missing revisions stay blank; a revision is not inferred from an arbitrary date in the document.

- **Policies:** review code, title, body and owner. Recognized policy tables can provide multiple policies. Otherwise a document-level policy candidate preserves the complete extracted text, including its sections. Missing metadata stays blank. A document containing multiple separate policies in an unrecognized layout should be split or prepared as a supported table rather than assuming automatic policy boundaries.

- **Requirements:** review reference and full text, then select an existing target framework in the current tenant. Recognized requirement tables and numbered normative clauses are supported. Framework existence and tenant ownership are checked again on import.

Candidate cards show a source page/line or DOCX block/table-row location and extraction warnings. PDF extraction preserves spacing for column mapping and wrapped descriptions. DOCX paragraphs and tables are read in document order. The extractor maps recognized column headers rather than assigning fields solely from their position. Unrecognized tables/aligned columns produce warnings. This is a deterministic source parser, not an authoritative standards interpretation or a model-derived accuracy score.

Repeated identifiers remain separate candidates with warnings. Review and consolidate the source content or select the appropriate occurrence; duplicate identifiers cannot be imported as separate records in the same scope. Text is not silently cut to fit destination fields. If a destination limit is exceeded, the import is rejected and the preview remains available for correction. Controls allow descriptions up to 10,000 characters; policies allow bodies up to 100,000 characters. A single extraction is limited to 1,000 candidates; split larger documents.

Select all and complete the import

1. Extract candidates for the intended record type and compare them with the source.

2. Fill missing required fields. For requirements, choose the existing target framework.

3. Choose **Select all** to include every extracted candidate, or select individual approval checkboxes. **Clear selection** removes every selection and cancels an older confirmation. Neither action imports data.

4. Check the selected-record count, then choose **Import selected reviewed requirements/controls/frameworks/policies**.

5. Check the count and target in the in-page confirmation, then choose **Confirm import**. Native browser confirmation dialogs are not used. **Cancel** leaves the database unchanged.

6. Wait for **Imported N reviewed records** and retain the import reference. For requirements, **Open target framework** displays the saved records.

7. If an error appears, retain its request reference for your administrator. A network interruption can leave the result uncertain: check the target before retrying. Existing controls are skipped without changes; check the receipt for created and skipped counts. Existing framework, policy and requirement conflicts reject the batch without partial writes.

Select all includes all candidates for the currently chosen record type. It does not remove required-field checks or automatically approve a standard as an organization policy. Controls remain planned and policies remain draft after a reviewed import.

Import behavior

Select only reviewed candidates and complete every required field. Import validates the full selection before creating records. Invalid new rows or an invalid framework prevent the batch from being imported. Existing control codes and later repeated control codes are recorded as skipped while new controls continue; original controls are not overwritten. Other record types retain their conflict safeguards. Conflicts do not discard the preview. Unexpected database errors roll back the transaction. Each successful record has an audit event identifying the reviewer and source file SHA-256.

Controls are created as **planned** and policies as **draft**, preserving their existing lifecycle and review requirements. Import does not approve or publish a policy and does not automatically link controls to frameworks. Tenant administrators and risk managers may extract/import; ordinary viewers may not. Existing CSV import and the reviewed-requirement endpoints remain available.

Scope and limitations

The upload limit is 8 MB, with at most 1,000 selected PDF pages and 2,000,000 source-text characters per review (generated layout spacing excluded), and bounded DOCX archive expansion. Image-only/scanned PDFs must undergo OCR before upload. Complex multi-column reading order, unusual tables, automatic Word numbering, embedded images and nested layouts may require source preparation or manual review. An empty candidate list is not a statement that a document contains no obligations.

Validate representative PDF and DOCX files against their original sections and table columns before importing. Source files and content are not sent to an external extraction or AI service.

ISO/IEC 27001:2022 PDFs

When the PDF's source content identifies ISO/IEC 27001:2022 and its Annex A structure, a dedicated profile uses content-stream text instead of layout-mode spacing. This avoids fragmented words from justified typography. It excludes cover/copyright material, contents/front matter, publisher page furniture and the bibliography from candidate records.

Choose **Frameworks** for the standard identifier and edition, **Requirements** for normative management clauses in sections 4-10, or **Controls** for Annex A table rows. Annex control references are namespaced as `A.5.1`, etc., so they cannot be confused with management clauses such as `5.1`. Wrapped titles, the table's `Control` column marker and split source references are handled separately. The detected standard is displayed in the preview. Counts describe what was extracted; the profile does not fabricate missing rows to meet a target.

Choosing **Policies** for this standard returns no organization policy and explains why: a standard sets policy requirements but does not contain your organization's approved policy text. Import its controls/requirements and draft your organization policies through the normal lifecycle.

For ISO/IEC 27001:2022, verify that the review separates the framework, management requirements and 93 Annex A controls (37 organizational, 8 people, 14 physical and 34 technological). Different editions, translations, scans and encodings require their own review. Missing ownership must be assigned explicitly. The standard is not an organization-approved policy. Licensed source documents are not included in this distribution.

Refresh the browser after an upgrade before re-extracting documents so the typed preview and detected-standard label are loaded. Previously imported records are not automatically deleted or replaced.

Import requirements when no framework exists

Requirements must belong to a target framework. If the target list is empty, enter the framework name and its actual revision / edition directly on Document Extraction and select **Create framework and use as target**. Creating a target does not import requirements. It preserves all extracted rows, edited text and selections. Choose **Import selected reviewed requirements**, review the target and count, then choose **Confirm import**. **Select all** approves every extracted row for the same review flow.

If you click Import without a target, the page explains the missing framework and focuses its name field. A failed framework creation leaves the preview intact and shows an error with its request reference. Framework creation and requirement import are separate audited transactions; the missing-target UI event is also logged without document contents.

Large PDF documents and page ranges

The old 250-page rejection is removed. Leave PDF start/end pages blank to process the complete document when it fits the per-review limits. Enter one-based inclusive page numbers to process a range (for example, 251 to 500). A blank start means page 1; a blank end means the final page. Page ranges apply only to PDF files. The result names the processed range and total page count; partial results explicitly warn that the remaining pages were not assessed. Candidate locations retain original PDF page numbers.

Review and import a range before extracting the next range, because a new extraction replaces the current unsaved preview. Use contiguous boundaries when numbered records continue across pages; inspect boundary records against the source before import. If a single record spans a boundary, include it in one reviewed range and avoid selecting duplicate references in the other. No automatic truncation or automatic import is performed. There is no fixed candidate or reviewed-record count ceiling. The 2-million source-text-character budget remains; reduce the range if the actual source exceeds it. The upload limit remains 8 MB, including when selecting a range from a large source PDF. HTTP errors retain their request ID in system and operations logs. Completion logs include processed start/end and total page count without source document text.

NIST SP 800-53 Revision 5 catalogue extraction

The Revision 5 PDF is recognized from its publication identity. The dedicated profile reads Chapter Three control statements in source order, including enhancement identifiers such as AC-2(1). It excludes the cover/front matter, explanatory Discussion and Related Controls sections, references and Appendix C summary tables. Withdrawn entries are omitted with an explicit count; they must not become planned active controls. Owners remain blank until you supply an accountable owner. The standard is not an organization policy: choose Controls or Requirements for its statements, or Frameworks for its publication identity.

For the 492-page NIST SP 800-53 Revision 5 PDF tested for this release, the review contains 1,007 active statements: 298 base controls and 709 enhancements. The profile omits 182 withdrawn entries. The full file can be reviewed in one extraction and one reviewed import; document review/import has no fixed record-count ceiling. CSV retains its separate file-size limit and has no fixed row-count ceiling. Review all extracted fields, set Default owner before extraction when appropriate, select records or Select all, then select Import and Confirm import. No controls are saved by extraction alone. Existing control codes are noted and skipped while new controls continue; existing policies, frameworks and requirements retain their existing conflict safeguards.

When using a page range, include the control heading and its complete statement. A range beginning inside an enhancement cannot establish its parent identifier; incomplete fragments are not invented as records. Original PDF page references remain attached to candidates.

PDF layout padding and extraction diagnostics

The PDF parser no longer emits artificial vertical blank-line padding. For generic tables, horizontal spacing still preserves column alignment, but generated spacing does not consume the source-text budget. A separate 64-million-character layout expansion guard bounds generated layout storage. The 8 MB upload, 1,000 selected-page and 2-million source-text limits remain explicit; no review is silently truncated.

Completion logs include plain-text, layout and content character counts together with the selected page range. Extraction rejections record these safe metrics and their HTTP status under the request ID; source document text and filenames are not logged. The NIST structural profile reads source-order text directly, avoiding layout padding and rotated publisher furniture. Existing imported data is preserved; refresh the browser and re-extract the original file after upgrading.

Choose the record type and verify the destination

Document Extraction now opens with **Choose records to create**. Choose the record type explicitly before extracting:

- **Controls** creates records in the **Controls** tab. Supply an accountable owner before import, using Default owner for missing owners when appropriate.

- **Requirements** creates records under **Frameworks -> selected framework -> Requirements**. Selecting a framework does not create control records. The Requirements navigation tab also provides viewing, adding and editing for the selected framework.

- **Frameworks** creates records in the **Frameworks** tab.

- **Policies** creates records in the **Policies** tab.

The destination is shown beside the type selector, in the extraction summary, and again before confirmation. For a detected standard, compare its identity with the chosen destination framework before proceeding. Intentional mappings to another framework remain possible; the product does not silently reclassify or move records.

After a successful import, the receipt identifies the destination and provides **Open Controls**, **Open imported requirements**, **Open Frameworks** or **Open Policies**. The requirements action opens the selected framework's requirements. Controls also offers **Import controls from a document**, which explicitly selects Controls in the extraction page. An empty Controls list explains that requirement imports are listed under Frameworks.

If a previous import completed but is absent from Controls, check its receipt or Frameworks requirements first. Existing requirements remain in their original framework. To create controls from a source standard, choose Controls and perform a separate reviewed import with explicit ownership. Do not assume that a successful Requirements import created controls.

Default accountable owner and duplicate control imports

Enter Default accountable owner before or after extraction. All rows that have no individual source/review owner inherit it. Changing the default updates inherited review fields; explicit individual owners remain unchanged. Clear an individual owner to restore inheritance. Changing the default cancels any pending confirmation so the final owner values must be reviewed again. The server applies the same fallback when an individual owner is omitted or blank in an API request. Policies use the same explicit-owner fallback; frameworks and requirements do not acquire invented ownership fields.

Existing controls are identified by exact, case-sensitive code within the current workspace. The review labels known existing codes. They are skipped without changing their owner, description, status, version or authorship. New controls continue importing; a repeated code inside one batch keeps its first occurrence and skips later occurrences. The server rechecks uniqueness at insertion, so an overlapping concurrent import is also treated as a skip. Completion reports created/skipped counts and lists each skipped code with its reason. These decisions are audited. A batch containing only existing controls succeeds with zero created records. Invalid new rows still prevent the new writes from being committed; duplicate skipping is not permission to save invalid data.

The same owner and duplicate-control rules apply to CSV import. Keep the owner column in the downloaded template; leave cells blank to inherit Default accountable owner, or specify an individual row owner. Changing the default after preview requires another preview. Existing policies/frameworks/requirements keep their original conflict behavior. CSV file-size limits remain; there is no fixed CSV row-count ceiling.

CJIS Security Policy version 6.1

CJIS is one published policy. Policies extraction retains all extracted text of the selected pages, including appendices when selected, as one draft policy candidate. For the authorized 473-page version 6.1 PDF, this is over one million characters; the body is retained rather than truncated into thousands of spurious table fragments. The accountable owner is supplied locally; publisher attribution does not become the importing organization's owner or approval.

Frameworks extraction returns the FBI CJIS Security Policy identity and version 6.1. Controls and Requirements extract 308 main-policy provisions: 292 control/control-enhancement statements and 16 normative mobile-device clauses. Source control references are preserved. Control codes use a CJIS- prefix to keep CJIS-specific provisions distinct from NIST controls with similar identifiers; requirement references remain scoped to the selected framework. Introductory context, discussion, examples and appendices remain available in the full policy text and are not invented as separate control requirements. This is source extraction, not an assertion of CJIS compliance or local implementation.

Requirements navigation

Open **Requirements** in the navigation menu. Choose the framework to view its requirements. Authorized users can choose **Add requirement**, enter its reference and text, and save it; existing entries can be edited through the same version checks used under Frameworks. If there is no framework, use **Create a framework** first. Read-only roles can view requirements. Requirements remain associated with their framework and can also still be viewed under Frameworks. The document-import receipt opens the selected framework directly in Requirements.